

Database Management Systems L5

Umass Boston
Summer 2023
Cristina Maier

Some slides are based on “Database Management System, 3rd ed., by Ramakrishnan and Gehrke

Topics

- ❖ Introduction to DBMS
- ❖ Relational Data Model
- ❖ *Relational Algebra*
- ❖ Conceptual Design: the Entity-Relationship Model
- ❖ **Structured Query Language (SQL)**
- ❖ Application Development (Java, Python)
- ❖ Schema Refinement and Normal Forms
- ❖ Database Security and Authorization
- ❖ Some NoSQL topics (If time permitted)

Example

Sailors

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.0
58	rusty	10	35.0
59	rusty	10	45.0
60	andy	10	60.0
61	mary	8	25.0

SELECT rating, MIN(age)
FROM Sailors
GROUP BY rating;

rating	MIN(age)
7	45.0
8	25.0
10	35.0

SELECT rating, MIN(age)
FROM Sailors
GROUP BY sname;

?

Example(cont.)

Sailors

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.0
58	rusty	10	35.0
59	rusty	10	45.0
60	andy	10	60.0
61	mary	8	25.0

SELECT rating, MIN(age)
FROM Sailors
GROUP BY rating;

rating	MIN(age)
7	45.0
8	25.0
10	35.0

SELECT rating, MIN(age)
FROM Sailors
GROUP BY sname;

Invalid query!
Why?

Example with HAVING

Sailors

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.0
58	rusty	10	35.0
59	rusty	10	45.0
60	andy	10	60.0
61	mary	8	25.0

```
SELECT rating, AVG(age)
FROM Sailors
GROUP BY rating;
```

rating	AVG(age)
7	45.0
8	40
10	46.66

```
SELECT rating, AVG(age)
FROM Sailors
GROUP BY rating
HAVING AVG(age) >= 45;
```

rating	AVG(age)
7	45.0
10	46.66

More Group qualification functions

- ❖ So far we have seen group qualification based on a property of the group
 - ❖ e.g.: aggregate function computed for each group
- ❖ Recent SQL standard allows group qualification based on a property of individual records
 - ❖ **EVERY(condition)**: True if condition holds for every group tuple
 - ❖ **ANY(condition)**: True if condition holds for any group tuple
 - ❖ *HAVING .. ANY .., HAVING ... EVERY ... are not currently supported by Oracle*

Example

Sailors3

sid	sname	rating	age	salary
22	dustin	7	45.0	20000
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

```
SELECT s.rating, MIN(s.age)
FROM Sailors3 s
GROUP BY s.rating;
```

rating	MIN(age)
7	45.0
8	25.0
10	35.0

```
SELECT s.rating, MIN(s.age)
FROM Sailors3 s
GROUP BY s.rating
HAVING EVERY (s.salary >= 20000);
```

rating	MIN(age)
7	45.0
8	25.0

Group with rating=10 is discarded because not all records from the that group have a salary >= 20000. Record with sid=38 does not satisfy the condition.

Next

- ❖ More Group By Exercises
- ❖ Nested Queries

Example 1

Find the min age for each rating. Include only data with
Sailors3 salaries greater than 25000.

sid	sname	rating	age	salary
22	dustin	7	45.0	20000
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

```
SELECT s.rating, MIN(s.age)
FROM Sailors3 s
WHERE s.salary > 25000
GROUP BY s.rating
```

WHERE salary >25000

sid	sname	rating	age	salary
22	dustin	7	45.0	20000
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

The GROUP by will be applied to

sid	sname	rating	age	salary
31	lubber	8	55.0	30000
59	rusty	10	45.0	40000
61	mary	8	25.0	30000

Example 1 (cont.)

Find the min age for each rating. Include only data with
Sailors3 salaries greater than 25000.

sid	sname	rating	age	salary
22	dustin	7	45.0	20000
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

```
SELECT s.rating, MIN(s.age)
FROM Sailors3 s
WHERE s.salary > 25000
GROUP BY s.rating
```

GROUP BY s.rating

rating	MIN(s.age)
8	25.0
10	45.0

Result is:

rating	MIN(s.age)
8	25.0
10	45.0

Example 2

Sailors3

Find the AVG salary for each rating
greater than 7.

Keep only the ones with an avg Salary <29000

sid	sname	rating	age	salary
22	dustin	7	45.0	20000
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

```
SELECT s.rating, AVG(s.salary)
FROM Sailors3 s
WHERE s.rating >7
GROUP BY s.rating
HAVING AVG(s.salary) < 29000;
```

Example 2

Find the AVG salary for each rating
greater than 7.

Keep only the ones with an avg Salary <29000

Sailors3

sid	sname	rating	age	salary
22	dustin	7	45.0	20000
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

```
SELECT s.rating, AVG(s.salary)
FROM Sailors3 s
WHERE s.rating > 7
GROUP BY s.rating
HAVING AVG(s.salary)<29000
```

WHERE s.rating >7

sid	sname	rating	age	salary
22	dustin	7	45.0	20000
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

The GROUP by will be applied to

sid	sname	rating	age	salary
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

Example 2 (cont.)

Find the AVG salary for each rating
greater than 7.

Keep only the ones with an avg Salary <29000

Sailors3

sid	sname	rating	age	salary
22	dustin	7	45.0	20000
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

```
SELECT s.rating, AVG(s.salary)
FROM Sailors3 s
WHERE s.rating > 7
GROUP BY s.rating
HAVING AVG(s.salary)<29000
```

The GROUP by will be applied to

sid	sname	rating	age	salary
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

GROUP BY s.rating

rating	AVG(salary)
8	30000
10	26666.66

Example 2 (cont.)

Find the AVG salary for each rating
greater than 7.

Sailors3

Keep only the ones with an avg Salary <29000

sid	sname	rating	age	salary
22	dustin	7	45.0	20000
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

```
SELECT s.rating, AVG(s.salary)
FROM Sailors3 s
WHERE s.rating > 7
GROUP BY s.rating
HAVING AVG(s.salary)<29000
```

HAVING AVG(s.salary)>29000

rating	AVG(salary)
8	30000
10	25666.66

Result is:

rating	AVG(salary)
10	26666.66

Example 3

(The aggregate from HAVING does not
need to be present in the target list)
Sailors3

sid	sname	rating	age	salary
22	dustin	7	45.0	20000
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

Find the AVG salary for each rating
greater than 7.

Keep only the ratings with an MIN Salary >20000

SQL Query?

Example 3

(The aggregate from HAVING does
not to be present in the target list)
Sailors3

sid	sname	rating	age	salary
22	dustin	7	45.0	20000
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

Find the AVG salary for each rating
greater than 7.

Keep only the ones with an MIN Salary >20000

```
SELECT s.rating, AVG(s.salary)
FROM Sailors3 s
WHERE s.rating > 7
GROUP BY s.rating
HAVING MIN(s.salary)>20000;
```

WHERE s.rating >7

sid	sname	rating	age	salary
22	dustin	7	45.0	20000
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

The GROUP by will be applied to

sid	sname	rating	age	salary
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

Example 3 (cont.)

Sailors3

sid	sname	rating	age	salary
22	dustin	7	45.0	20000
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

Find the AVG salary for each rating
greater than 7.
Keep only the ones with an MIN Salary >20000.

```
SELECT s.rating, AVG(s.salary)
FROM Sailors3 s
WHERE s.rating > 7
GROUP BY s.rating
HAVING MIN(s.salary)>20000;
```

The GROUP by will be applied to

sid	sname	rating	age	salary
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

rating	MIN(salary)	AVG(salary)
8	30000	30000
10	15000	25666.66

Example 3 (cont.)

Sailors3

sid	sname	rating	age	salary
22	dustin	7	45.0	20000
31	lubber	8	55.0	30000
58	rusty	10	35.0	15000
59	rusty	10	45.0	40000
60	andy	10	60.0	25000
61	mary	8	25.0	30000

Find the AVG salary for each rating
greater than 7.
Keep only the ones with an MIN Salary >20000.

```
SELECT s.rating, AVG(s.salary)
FROM Sailors3 s
WHERE s.rating > 7
GROUP BY s.rating
HAVING MIN(s.salary)>20000;
```

HAVING MIN(salary)> 20000

rating	MIN(salary)	AVG(salary)
8	30000	30000
10	15000	25666.66

Result:
(Not that it has the attributes
and aggregates from target-list.)

rating	AVG(salary)
8	30000

SQL

- ❖ Create, alter, delete tables
- ❖ Insert Statement
- ❖ Basic Select Statement Structure
- ❖ LIKE, AS keywords
- ❖ Dates, Text case sensitivity
- ❖ Count
- ❖ Set Operations
- ❖ Aggregates
- ❖ Nested Queries

Nested Queries

- ❖ An SQL Query can be used to help evaluate another query
- ❖ E.g.: a condition might need to be evaluated on a computed relation, not on one readily available
- ❖ Multiple levels of nesting are possible
- ❖ Semantics are similar to those of nest loops

Connecting queries and sub-queries

- ❖ The result of one query can be:

- ❖ A relation

- ❖ E.g.: `SELECT s.sname, s.rating FROM Sailors s;`

- ❖ A scalar value (1×1 table)

- ❖ In this case, this result can be used in locations where a constant can be placed

- ❖ E.g.: `SELECT AVG(s.rating) from Sailors s;`

- ❖ Where do subqueries appear?

- ❖ Most often in the **WHERE** clause of a parent query

- ❖ In **FROM** clause follow by range variable

- ❖ In **HAVING** clause

Examples Database Instance

sailors

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.0
58	rusty	10	35.0
59	rusty	10	45.0

reserves

sid	bid	day
22	101	10/10/22
58	101	10/11/22
22	102	10/20/22

boats

bid	name	color
101	interlake	red
102	clipper	green

Subqueries that return a scalar

- ❖ Also referred to as subqueries that return a constant
- ❖ *If the subquery returns zero or more values, an error will occur. Very fragile!*
- ❖ Example: Select all sailors who have the smallest rating

```
SELECT * from sailors s
```

```
WHERE s.rating = (SELECT min(s2.rating)  
                  FROM sailors s2);
```

Conditions involving relations

- ❖ Test that a relation satisfies some condition:
 - ❖ `WHERE EXISTS (SELECT ...)` returns True if the subquery result is not empty
 - ❖ `WHERE UNIQUE (SELECT ...)` returns True if the subquery result has no duplicates
- ❖ E.g: Select all sailors for which there is a reservation for boat 101.

```
SELECT *  
FROM Sailors s  
WHERE EXISTS (SELECT *  
              FROM Reserves r  
              WHERE r.bid=101 and s.sid=r.sid);
```
- ❖ Subquery is **correlated** with parent query

Conditions involving relations and tuples

- ❖ Typically have some sort of set-operations semantics
 - ❖ WHERE field **IN** (SELECT ...)
 - ❖ WHERE field op **ANY** (SELECT ...)
 - ❖ WHERE field op **ALL** (SELECT ...)
- ❖ E.g.: select all sailors who have at least one reservation
SELECT * from Sailors s
WHERE s.sid **IN** (SELECT r.sid
FROM Reserves r);

Conditions involving relations and tuples

- ❖ Typically have some sort of set-operations semantics
 - ❖ WHERE field **IN** (SELECT ...)
 - ❖ WHERE field op **ANY** (SELECT ...)
 - ❖ WHERE field op **ALL** (SELECT ...)
- ❖ E.g.: select all sailors who have no reservation
SELECT * from Sailors s
WHERE s.sid **NOT IN** (SELECT r.sid
FROM Reserves r);

Conditions involving relations and tuples

- ❖ Typically have some sort of set-operations semantics
 - ❖ WHERE field **IN** (SELECT ...)
 - ❖ WHERE field op **ANY** (SELECT ...)
 - ❖ WHERE field op **ALL** (SELECT ...)
- ❖ E.g.: select all sailors whose rating is higher than the rating of any sailor who reserved boat 101

```
SELECT * from Sailors s
```

```
WHERE s.rating > ANY ( SELECT s2.rating  
                        FROM Sailors s2,Reserves r  
                        WHERE s2.sid=r.sid and r.bid=101);
```

Conditions involving relations and tuples

- ❖ Typically have some sort of set-operations semantics
 - ❖ WHERE field **IN** (SELECT ...)
 - ❖ WHERE field op **ANY** (SELECT ...)
 - ❖ WHERE field op **ALL** (SELECT ...)
- ❖ E.g.: select all sailors with the oldest age
SELECT * from Sailors s
WHERE s.age **>= ALL** (SELECT s2.age
FROM Sailors s2);

Same query could be written with aggregates

❖ E.g.: select all sailors with the oldest age

```
SELECT * from Sailors s
```

```
WHERE s.age = (SELECT max(s2.age)  
              FROM Sailors s2);
```

Oracle Practice Session

- ❖ Connect to SqlPlus
- ❖ PracticeSessionSQL3.sql

SQL

- ❖ Create, alter, delete tables
- ❖ Insert Statement
- ❖ Basic Select Statement Structure
- ❖ LIKE, AS keywords
- ❖ Dates, Text case sensitivity
- ❖ Count
- ❖ Set Operations
- ❖ Aggregates
- ❖ Nested Queries
- ❖ SQL Division

Examples Database Instance

sailors

sid	sname	rating	age
22	dustin	7	45.0
31	lubber	8	55.0
58	rusty	10	35.0
59	rusty	10	45.0

reserves

sid	bid	day
22	101	10/10/22
58	101	10/11/22
22	102	10/20/22

boats

bid	name	color
101	interlake	red
102	clipper	green

Recap.: Conditions involving relations

- ❖ Test that a relation satisfies some condition:
 - ❖ `WHERE EXISTS (SELECT ...)` returns True if the subquery result is not empty
 - ❖ `WHERE UNIQUE (SELECT ...)` returns True if the subquery result has no duplicates
- ❖ E.g: Select all sailors for which there is a reservation for boat 101.

```
SELECT *  
FROM Sailors s  
WHERE EXISTS (SELECT *  
              FROM Reserves r  
              WHERE r.bid=101 and s.sid=r.sid);
```
- ❖ Subquery is **correlated** with parent query

Recap.: Conditions involving relations and tuples

- ❖ Typically have some sort of set-operations semantics
 - ❖ WHERE field **IN** (SELECT ...)
 - ❖ WHERE field op **ANY** (SELECT ...)
 - ❖ WHERE field op **ALL** (SELECT ...)
- ❖ E.g.: select all sailors who have at least one reservation
SELECT * from Sailors s
WHERE s.sid **IN** (SELECT r.sid
FROM Reserves r);

Subqueries in the FROM clause

- ❖ Find the names of sailors who reserved boat interlake. Return the sailors names and the boat name.

```
SELECT SQ.sname, SQ.bname
FROM (SELECT S.sname, B.name as bname
      FROM Sailors S, Boats B, Reserves R
      WHERE S.sid=R.sid and R.bid=B.bid) SQ
WHERE SQ.bname='interlake';
```

❖

HAVING with Subquery

- ❖ Find the age of the youngest sailor older than 18, for each rating with at least two sailors.

```
SELECT S.rating, min(S.age) FROM Sailors S
WHERE S.age >= 18
GROUP BY S.rating
HAVING 1 < (SELECT COUNT(*) FROM Sailors S2
            WHERE S.rating=S2.rating);
```

- ❖ NOTE: HAVING clause always executes after WHERE clause

❖

Example

Incorrect solution

- ❖ sailors(sid:int, sname: string, rating:int, age:, salary:real)
- ❖ boats(bid:int, name:string, color:string, manufacturer:string, prod_year:date)
- ❖ reserves (sid:int, bid:int, day:date)
- ❖ Find the sailors' ids who reserved both red and green boats
- ❖ The solution given below is incorrect !!!

```
SELECT s.sid
```

```
FROM sailors s, reserves r, boats b
```

```
WHERE s.sid = r.sid AND b.bid =r.bid AND
```

```
    b.color='red' AND b.color='green';
```

Rewriting INTERSECT queries using IN

❖ Find the sailors' ids who reserved both red and green boats

❖ Correct Solution with INTERSECT:

```
(SELECT r.sid  
FROM Reserves r, Boats b  
WHERE r.bid=b.bid AND b.color='red')
```

INTERSECT

```
(SELECT r.sid  
FROM Reserves r, Boats b  
WHERE r.bid = b.bid AND b.color='green');
```

Rewriting INTERSECT queries using IN

❖ Find the sailors' ids who reserved both red and green boats

❖ Solution without INTERSECT:

```
SELECT DISTINCT r.sid
```

```
FROM Reserves r, Boats b
```

```
WHERE r.bid=b.bid AND
```

```
    b.color='red' AND r.sid IN ( SELECT r2.sid
```

```
                                FROM Reserves r2, Boats b2
```

```
                                WHERE r2.bid=b2.bid AND b2.color='green');
```

Another Example

- ❖ sailors(sid:int, sname: string, rating:int, age:, salary:real)
- ❖ boats(bid:int, name:string, color:string, manufacturer:string, prod_year:date)
- ❖ reserves (sid:int, bid:int, day:date)
- ❖ Find the boat ids and names that are booked at least 3 times

SELECT b.bid, b.name

FROM boats b

Where 3 <= (select count(*)

FROM reserves r2

Where **r2.bid=b.bid**);

Another Example

- ❖ sailors(sid:int, sname: string, rating:int, age:, salary:real)
- ❖ boats(bid:int, name:string, color:string, manufacturer:string, prod_year:date)
- ❖ reserves (sid:int, bid:int, day:date)
- ❖ Find the boat ids and names that are booked by at least 3 unique sailors

```
SELECT b.bid, b.name
```

```
FROM boats b
```

```
Where 3 <= (select count(DISTINCT r2.sid)
```

```
FROM reserves r2
```

```
Where r2.bid=b.bid);
```

Nested Queries Summary

- ❖ Nested queries returning a constant
 - ❖ Typically the constant is compared with other values in the WHERE clause
 - ❖ ... WHERE field= (SELECT x FROM...)...
- ❖ Nested queries returning a relation
 - ❖ In WHERE Clause
 - ❖ ...WHERE (EXISTS/UNIQUE) (SELECT...)...
 - ❖ ... WHERE field IN (SELECT col FROM)...
 - ❖ ...WHERE field op ANY|ALL (SELECT col FROM...)....
 - ❖ In FROM clause, followed by range variable
 - ❖ E.g: ..FROM Sailors, (SELECT bid FROM Boats..) BIDS
 - ❖ In HAVING clause
 - ❖ ... HAVING field/const op (SELECT)

Questions?