

Database Management Systems L6

Umass Boston
Summer 2023
Cristina Maier

Some slides are based on “Database Management System, 3rd ed., by Ramakrishnan and Gehrke

Topics

- ❖ Introduction to DBMS
- ❖ Relational Data Model
- ❖ *Relational Algebra*
- ❖ Conceptual Design: the Entity-Relationship Model
- ❖ **Structured Query Language (SQL)**
- ❖ Application Development (Java, Python)
- ❖ Schema Refinement and Normal Forms
- ❖ Database Security and Authorization
- ❖ Some NoSQL topics (If time permitted)

Connecting queries and sub-queries (recap.)

- ❖ The result of one query can be:
 - ❖ A relation
 - ❖ E.g.: `SELECT s.sname, s.rating FROM Sailors s;`
 - ❖ A scalar value (1×1 table)
 - ❖ In this case, this result can be used in locations where a constant can be placed
 - ❖ E.g.: `SELECT AVG(s.rating) from Sailors s;`
- ❖ Where do subqueries appear?
 - ❖ Most often in the WHERE clause of a parent query
 - ❖ In **FROM** clause follow by range variable
 - ❖ In **HAVING** clause

Nested Queries Summary (recap.)

- ❖ Nested queries returning a constant
 - ❖ Typically the constant is compared with other values in the WHERE clause
 - ❖ ... WHERE field= (SELECT x FROM...)...
- ❖ Nested queries returning a relation
 - ❖ In WHERE Clause
 - ❖ ...WHERE (EXISTS/UNIQUE) (SELECT...)...
 - ❖ ... WHERE field IN (SELECT col FROM)...
 - ❖ ...WHERE field op ANY|ALL (SELECT col FROM...)....
 - ❖ In FROM clause, followed by range variable
 - ❖ E.g: ..FROM Sailors, (SELECT bid FROM Boats..) BIDS
 - ❖ In HAVING clause
 - ❖ ... HAVING field/const op (SELECT)

Exercise 1

- ❖ sailors(sid:int, sname: string, rating:int, age:, salary:real)
- ❖ boats(bid:int, name:string, color:string, manufacturer:string, prod_year:date)
- ❖ reserves (sid:int, bid:int, day:date)
- ❖ Find the id and name of the oldest boats

```
SELECT b.bid, b.name
```

```
FROM boats b
```

```
WHERE b.prod_year = ( SELECT MIN(b2.prod_year)  
                     FROM boats b2);
```

Exercise 2

- ❖ sailors(sid:int, sname: string, rating:int, age:, salary:real)
- ❖ boats(bid:int, name:string, color:string, manufacturer:string, prod_year:date)
- ❖ reserves (sid:int, bid:int, day:date)
- ❖ Find the sailors who DO NOT have the highest rating. Extract the sailors id and name.

Exercise 2 - Solution #1

- ❖ sailors(sid:int, sname: string, rating:int, age:, salary:real)
- ❖ boats(bid:int, name:string, color:string, manufacturer:string, prod_year:date)
- ❖ reserves (sid:int, bid:int, day:date)
- ❖ Find the sailors who DO NOT have the highest rating. Extract the sailors id and name.

```
SELECT s.sid, s.sname, s.rating, s.age, s.salary
FROM sailors s
WHERE s.rating != ( SELECT MAX(s2.rating)
                    FROM sailors s2);
```

Exercise 2 - solution #2

(using ANY)

- ❖ sailors(sid:int, sname: string, rating:int, age:, salary:real)
- ❖ boats(bid:int, name:string, color:string, manufacturer:string, prod_year:date)
- ❖ reserves (sid:int, bid:int, day:date)
- ❖ Find the sailors who DO NOT have the highest rating. Extract the sailors id and name.

```
SELECT s.sid, s.sname, s.rating, s.age, s.salary
FROM sailors s
WHERE s.rating < ANY ( SELECT s2.rating
                       FROM sailors s2);
```


Exercise 2 - solution #3

```
SELECT s.sid, s.sname, s.rating, s.age, s.salary
FROM sailors s
WHERE EXISTS ( SELECT *
                FROM sailors s2
                WHERE s2.pub_year >
s.pub_year);
```

SQL Division

- ❖ Remember the division from Relational Algebra
- ❖ SQL does not implement the division (/) operator

Division (recap)

$$R/S$$

- ❖ Also noted as $R \div S$
- ❖ Useful for expressing certain kinds of queries, such as “Find the names of all sailors who have reserved all boats.”
- ❖ It does not have the same importance as the other operators because it is not needed as often, but it is extremely important for certain queries
 - ❖ SQL does not have a specific construct for it. When needed, it is done using a combination of other constructs

Division (recap.)

- ❖ $A(x:\text{type1}, y:\text{type2})$
- ❖ $B(y:\text{type2})$
- ❖ A/B is the set of all values x (in the form of unary tuples) such that for every y value in B , there is a $\langle x, y \rangle$ tuple in A
- ❖ For each x value in the first column of A , consider the set of values that appear in the y column of A associated with that value of x . If this set contains all values y in B , then x is in the result for A/B
- ❖ A value x from A is disqualified from the result if by attaching a value y from B , we have a tuple $\langle x, y \rangle$ that is not present in A
- ❖ $A/B = \pi_x A - \pi_x((\pi_x(A) \times B) - A)$
- ❖ To understand the division operation in full generality, we have to consider the case when both x and y are replaced by a set of attributes

Division Example (recap.)

R

sid	bid
12	101
10	102
12	102
11	101
10	101

B

bid
101
102

Find the ids of sailors who have reserved all boats

R/B

sid
12
10

Division in SQL

- ❖ SQL does not implement the division operator
- ❖ One of the most subtle query
- ❖ Solutions follow specific patterns

Division in SQL

- ❖ Two ways of writing it
 - ❖ 1) using the set-difference operator (EXCEPT, MINUS)
 - ❖ 2 levels of nested queries
 - ❖ 2) without set-difference (EXCEPT, MINUS)
 - ❖ 3 levels of nested queries

Schema

- ❖ sailors(sid:int, sname: string, rating:int, age:, salary:real)
- ❖ boats(bid:int, name:string, color:string, manufacturer:string, prod_year:date)
- ❖ reserves (sid:int, bid:int, day:date)

Solution with set-difference (EXCEPT or MINUS)

- ❖ Find the sailors who reserved all boats
- ❖ This is equivalent to *'find all sailors for whom there is no boat they did not reserved'*

```
SELECT s.sid, s.sname
FROM sailors s
WHERE NOT EXISTS (
    ( SELECT b.bid FROM boats b)
    MINUS
    (SELECT r.bid FROM reserves r
    WHERE r.sid=s.sid)
);
```

Solution without set-difference (EXCEPT or MINUS)

- ❖ Find the sailors who reserved all boats
- ❖ SELECT *a sailor for whom*
- ❖ *there is no boat*
- ❖ *for which there is no reservation made by that sailor*

```
SELECT s.sid, s.sname
FROM sailors s
WHERE NOT EXISTS (
    SELECT b.bid FROM boats b
    WHERE NOT EXISTS (
        SELECT * FROM reserves r
        WHERE r.sid=s.sid AND r.bid=b.bid
    )
);
```

Example 2

- ❖ `Movies(movie_id: int, title: string, year: int, studio: string)`
 - ❖ `Actors(actor_id: int, name: string, age: int)`
 - ❖ `PlaysIn(actor_id: int, movie_id: int, character: string)`
 - ❖ Find the id and name of actors who played in all movies
- ?

Example 2 - solution #1

- ❖ Movies(movie_id: int, title: string, year: int, studio: string)
- ❖ Actors(actor_id: int, name: string, age: int)
- ❖ PlaysIn(actor_id: int, movie_id: int, character: string)
- ❖ Find the id and name of actors who played in all movies

```
SELECT a.actor_id, a.name
```

```
FROM Actors a
```

```
WHERE NOT EXISTS(
```

```
    (SELECT m.movie_id FROM Movies m)
```

```
    MINUS
```

```
    (SELECT p.movie_id FROM PlaysIn p
```

```
    WHERE p.actor_id=a.actor_id)
```

```
);
```

Example 2 - solution #2

- ❖ Movies(movie_id: int, title: string, year: int, studio: string)
- ❖ Actors(actor_id: int, name: string, age: int)
- ❖ PlaysIn(actor_id: int, movie_id: int, character: string)
- ❖ Find the id and name of actors who played in all movies

```
SELECT a.actor_id, a.name
```

```
FROM Actors a
```

```
WHERE NOT EXISTS (
```

```
    SELECT m.movie_id FROM Movies m
```

```
    WHERE NOT EXISTS (
```

```
        SELECT * from PlaysIn p
```

```
        WHERE p.actor_id=a.actor_id AND p.movie_id=m.movie_id
```

```
    )
```

```
);
```

Example 3

- ❖ `Movies(movie_id: int, title: string, year: int, studio: string)`
- ❖ `Actors(actor_id: int, name: string, age: int)`
- ❖ `PlaysIn(actor_id: int, movie_id: int, character: string)`
- ❖ Find the id and name of actors who played in all movies produced by 'Universal' studio.

?

Example 3 - solution #1

- ❖ Movies(movie_id: int, title: string, year: int, studio: string)
- ❖ Actors(actor_id: int, name: string, age: int)
- ❖ PlaysIn(actor_id: int, movie_id: int, character: string)
- ❖ Find the id and name of actors who played in all movies produced by 'Universal' studio.

```
SELECT a.actor_id, a.name
```

```
FROM Actors a
```

```
WHERE NOT EXISTS(
```

```
  (SELECT m.movie_id FROM Movies m
```

```
   WHERE m.studio='Universal')
```

```
  MINUS
```

```
  (SELECT p.movie_id FROM PlaysIn p
```

```
   WHERE p.actor_id=a.actor_id)
```

```
);
```

Example 3 - solution #2

- ❖ Movies(movie_id: int, title: string, year: int, studio: string)
- ❖ Actors(actor_id: int, name: string, age :int)
- ❖ PlaysIn(actor_id:int, movie_id: int, character: string)
- ❖ Find the id and name of actors who played in all movies produced by 'Universal' studio.

```
SELECT a.actor_id, a.name
```

```
FROM Actors a
```

```
WHERE NOT EXISTS (
```

```
    SELECT m.movie_id FROM Movies m
```

```
    WHERE m.studio='Universal' AND NOT EXISTS (
```

```
        SELECT * from PlaysIn p
```

```
        WHERE p.actor_id=a.actor_id AND p.movie_id=m.movie_id
```

```
    )
```

```
);
```


Example 4

- ❖ `Movies(movie_id: int, title: string, year: int, studio: string)`
- ❖ `Actors(actor_id: int, name: string, age: real)`
- ❖ `PlaysIn(actor_id:int, movie_id: int, character: string)`
- ❖ Find the average age of actors for each movie in which at least 4 actors play.
- ❖ ?

Example 4 - solution

- ❖ Movies(movie_id: int, title: string, year: int, studio: string)
- ❖ Actors(actor_id: int, name: string, age: real, age:int)
- ❖ PlaysIn(actor_id:int, movie_id: int, character: string)
- ❖ Find the average age of actors for each movie in which at least 4 actors play.

```
SELECT m.movie_id, AVG(a.age)
FROM Movies m, Actors a, PlaysIn p
WHERE m.movie_id=p.movie_id AND p.actor_id=a.actor_id
GROUP BY m.movie_id
HAVING COUNT(*) >=4;
```

Example 4 - solution #2

- ❖ Movies(movie_id: int, title: string, year: int, studio: string)
- ❖ Actors(actor_id: int, name: string, age: real)
- ❖ PlaysIn(actor_id:int, movie_id: int, character: string)
- ❖ Find the average age of actors for each movie in which at least 4 actors play.

```
SELECT p.movie_id, avg(a.age)
FROM Actors a, PlaysIn p
WHERE a.actor_id=p.actor_id
GROUP BY p.movie_id
HAVING 4 <= ( SELECT COUNT(p2.actor_id)
              FROM PlaysIn p2
              WHERE p2.movie_id=p.movie_id
            );
```

Practice queries

- ❖ Please practice at home with queries from PracticeSessionSQL4.sql

SQL

- ❖ Create, alter, delete tables
- ❖ Insert Statement
- ❖ Basic Select Statement Structure
- ❖ LIKE, AS keywords
- ❖ Dates, Text case sensitivity
- ❖ Count
- ❖ Set Operations
- ❖ Aggregates
- ❖ Nested Queries
- ❖ SQL Division
- ❖ NULL Constraints

Example Schema

- ❖ Movies(movie_id: int, title: string, year: int, studio: string)
- ❖ Actors(actor_id: int, name: string, age: int)
- ❖ PlaysIn(actor_id:int, movie_id: int, character: string)

NULL Constraints

- ❖ It has to do with missing values from columns
- ❖ Columns that do not have the NOT NULL constraint can contain null values
- ❖ PRIMARY KEYS have NOT NULL constraints !!!
- ❖ If you try to insert a record and do not provide values for primary key columns or for NOT NULL columns, you will get an error
- ❖ We can use conditions with **IS NOT NULL** to filter out all records that have a null value for that column
- ❖ We can use conditions with **IS NULL** to filter out all records that have a non-null value for that column

Example - Table Definition

- ❖ CREATE TABLE Movies(
 movie_id NUMBER(9) PRIMARY KEY,
 title VARCHAR(2),
 year int,
 studio VARCHAR(20)
);
- ❖ Column movie_id cannot contain null values, because it is a primary key. By default primary keys have NOT NULL constraints!
 - ❖ When inserting a record, movie_id must be present. If not present, we will get an error
- ❖ Columns title, year and studio can have null values. There is no NOT NULL constraint
 - ❖ When inserting a record, it is ok if we do not provide values for title, year or/and studio. If not present, they will contain null

Example 2 - Table Definition

- ❖ CREATE TABLE Movies2(
 movie_id NUMBER(9) PRIMARY KEY,
 title VARCHAR(2),
 year int,
 studio VARCHAR(20) NOT NULL
);
- ❖ Column movie_id cannot contain null values, because it is a primary key. By default primary keys have NOT NULL constraints!
- ❖ Column studio cannot contain null values because it has the NOT NULL constraint
- ❖ Columns title and year can have null values. There is no NOT NULL constraint

Example IS NOT NULL

- ❖ Movies(movie_id: int, title: string, year: int, studio: string)
- ❖ Actors(actor_id: int, name: string, age :int)
- ❖ PlaysIn(actor_id:int, movie_id: int, character: string)
- ❖ Find all movies that have a studio (i.r. column studio has a value)

SELECT *

FROM Movies

WHERE studio IS NOT NULL;

Example IS NULL

- ❖ Movies(movie_id: int, title: string, year: int, studio: string)
- ❖ Actors(actor_id: int, name: string, age :int)
- ❖ PlaysIn(actor_id:int, movie_id: int, character: string)
- ❖ Find all movies that DO NOT have a studio (i.r. column studio is null)

SELECT *

FROM Movies

WHERE studio IS NULL;

Problems with NULL Values

- ❖ It complicates things
- ❖ E.g.:
- ❖ **WHERE rating >7**
- ❖ What if rating allows NULLs?
- ❖ For NULL values, how `NULL >7` evaluate?
- ❖ Expressions can have **three logic values**
 - ❖ **True**
 - ❖ **False**
 - ❖ **Unknown**
- ❖ For a record that does have rating Null, the Condition WHERE rating >7 will evaluate to Unknown

Problems with NULL Values

- ❖ WHERE condition
 - ❖ Returns only the records for where condition is evaluated to True !!!
- ❖ Unknown AND True evaluates to Unknown
- ❖ Unknown OR False evaluates to Unknown
- ❖ Unknown OR True evaluates to True
- ❖ Unknown = Unknown evaluates to Unknown

Null Values and Aggregates

- ❖ COUNT(*) counts all records (regardless on whether some attributes are NULL)
- ❖ COUNT(attribute) only counts the records in which this attribute is not null
- ❖ All other aggregates skips NULL values if the aggregate is on a field that is NULL (e.g. SUM(attribute) only adds up the values where attribute is not null)
 - ❖ If all values of the attribute from the aggregate are null the aggregate will return NULL, except the case when the aggregate is COUNT. If COUNT, it will return 0.

Example

- ❖ `SELECT COUNT(*) FROM Movies;`
- ❖ `SELECT COUNT(studio) FROM Movies;`
- ❖ Do not return the same result if studio contains NULL values
- ❖ `SELECT COUNT(*) FROM Movies m;`
- ❖ `SELECT COUNT(*) FROM Movies m
WHERE m.year < 2010 or m.year >= 2010;`
- ❖ Do not return the same result if year contains NULL

Null Values And Duplicates

- ❖ When comparing two NULL values, the result is unknown
- ❖ There is an **anomaly** when checking for **duplicates**
 - ❖ In this case NULL values are considered equal
 - ❖ **DISTINCT, UNIQUE** keywords
 - ❖ If multiple rows have equal values in all non-null columns, then only one row is returned in the result

Group By

- ❖ Group by columnY
- ❖ If columnY contains NULL values, then a common group for NULL is created

Oracle Session

❖ Please practice at home with:

❖ PracticeSessionSQL4.sql

❖ PracticeSessionSQL5.sql

Questions?